



AN APPLIED MATERIALS
COMPANY

NEMA| GFX API

Application Note - Debugging

Version v.1.0

Part Number: D-DBG-p

June 25, 2025

Disclaimer

This document is written in good faith with the intent to assist the readers in the use of the product. Circuit diagrams and other information relating to Think Silicon S.A products are included as a means of illustrating typical applications. Although the information has been checked and is believed to be accurate, no responsibility is assumed for inaccuracies. Information contained in this document is subject to continuous improvements and developments.

Think Silicon S.A products are not designed, intended, authorized or warranted for use in any life support or other application where product failure could cause or contribute to personal injury or severe property damage. Any and all such uses without prior written approval of Think Silicon S.A. will be fully at the risk of the customer.

Think Silicon S.A. disclaims and excludes any and all warranties, including without limitation any and all implied warranties of merchantability, fitness for a particular purpose, title, and infringement and the like, and any and all warranties arising from any course or dealing or usage of trade.

This document may not be copied, reproduced, or transmitted to others in any manner. Nor may any use of information in this document be made, except for the specific purposes for which it is transmitted to the recipient, without the prior written consent of Think Silicon S.A. This specification is subject to change at anytime without notice.

Think Silicon S.A. is not responsible for any errors contained herein. In no event shall Think Silicon S.A. be liable for any direct, indirect, incidental, special, punitive, or consequential damages; or for loss of data, profits, savings or revenues of any kind; regardless of the form of action, whether based on contract; tort; negligence of Think Silicon S.A or others; strict liability; breach of warranty; or otherwise; whether or not any remedy of buyers is held to have failed of its essential purpose, and whether or not Think Silicon S.A. has been advised of the possibility of such damages.

COPYRIGHT NOTICE

NO PART OF THIS SPECIFICATION MAY BE REPRODUCED IN ANY FORM OR MEANS, WITHOUT THE PRIOR WRITTEN CONSENT OF THINK SILICON S.A.

Questions or comments may be directed to:

Think Silicon S.A

Suite B8

Patras Science Park

Rion Achaias 26504, Greece

web: <http://www.think-silicon.com>

email: info@think-silicon.com

Tel: +30 2610 911543

Fax: +30 2610 911544

Contents

Preface.....	4
About this Manual.....	4
Target Audience.....	4
Related Documents.....	4
Generic Workflow.....	5
Nothing is drawn.....	5
Frame is different than expected.....	6
Memory Leaks.....	7
Performance Issues.....	7
Breakpoints.....	9
GPU Status.....	11

1 Preface

1.1 About this Manual

This manual provides the necessary guidelines in order to perform debugging on applications that make use of NEMA| GFX.

1.2 Target Audience

The current document is intended to be used by engineers and developers who develop graphics applications using NEMA| GFX and therefore they are familiar with NEMA| GFX up to a certain degree.

1.3 Related Documents

The following documents are considered relevant for using the full spectrum of features of the NEMA| GFX Library:

- NEMA| GFX User Manual
- NEMA| pico Graphics Processing Unit Software User Manual
- NEMA| pico Graphics Processing Unit Hardware User Manual

2 Generic Workflow

Potential bugs that can be found in NEMA| GFX based applications can be divided into the following categories:

1. Nothing is drawn (the display is not updated)
2. Frame is different than expected (the drawn frame is not the expected one)
3. Memory leaks (allocated memory that never gets deallocated)
4. Performance issues (the performance is lower than expected)

Resolving each case can be performed by applying a set of actions that attempt to identify and resolve the target malfunction. These actions can be classified as:

- **[M]**: Mandatory
- **[A]**: Advisory
- **[N]**: Note

It must be noted that it is impossible to resolve the bug without performing the actions marked as Mandatory. Advisory actions refer to actions that have high potential of resolving the bug, provided that the mandatory actions have been performed. Finally, actions marked as Note, are generic advices that apply to each category. Although they do not comprise solid solutions, it is highly recommended that the users take them into consideration in order to minimize the chances of inserting bugs within their applications.

When creating a NEMA| GFX based application, it needs to be taken into account that the workflow follows a specific order. For instance, memory allocation precedes memory deallocation. Having this in mind, the following paragraphs present the actions that need to be applied, in the right order (from the beginning till the termination of the application).

2.1 Nothing is drawn

1. **[N]** A DC typically serves as a client and the GPU as a server. The GPU draws the FB and the DC reads it and provides the appropriate signaling to drive a display.
2. **[M]** `nema_init()` is the first NEMA| GFX function that is called.
3. **[M]** `nema_init()` should return with no errors (return 0)
4. **[M]** A FB should be allocated to an appropriate memory region. FB should be able to be written by the GPU and read by the DC.
5. **[N]** `nema_create_buffer()` can be used to allocate a FB.
6. **[N]** Multiple FBs (double/triple buffering) can be used to avoid visual artifacts, in which case synchronisation between DC and GPU must be taken care of.
7. **[N]** All NemaGFX functions that define the rendering pipeline (e.g. anything relative to setting blending mode, textures, fill/blit operations) operate on the bound CL. Nothing is drawn unless the CL is submitted for execution (`nema_cl_submit()`).
8. **[M]** A CL must be created by calling `nema_cl_create()`.

9. **[M]** All GPU context should be defined by the current CL, i.e. a CL must be agnostic of the current GPU state.
10. **[N]** A valid CL should contain at least the following: Clipping Rectangle, destination TEX (FB), blending mode, source TEX (in case of blitting), drawing operation.
11. **[A]** After the `nema_init()`, a CL should be bound at all times (`nema_cl_bind()`).
12. **[N]** Calling `nema_cl_bind()` implicitly unbinds the previously bound CL.
13. **[A]** Rewind CL before using it (`nema_cl_rewind()`).
14. **[M]** A destination surface must be bound before defining rendering commands (`nema_bind_dst()`). In most cases, this destination surface should correspond to an allocated FB.
15. **[M]** Clipping rectangle must be set within the boundaries of the FB (`nema_set_clip(x,y,w,h)`), and must fulfill the following conditions:
 - a. $x \geq 0$ and $y \geq 0$
 - b. $x+w \leq \text{fb.RESX}$ and $y+h \leq \text{fb.RESY}$
16. **[M]** A blending mode must be set before doing any drawing operations (blit or fill) (`nema_set_blend_blit()` / `nema_set_blend_fill()`).
17. **[M]** When blitting, a Source TEX containing an allocated texture must be bound (`nema_bind_src()`) before calling the corresponding rendering commands.
18. **[N]** PRIMs are drawn using the respective function calls (e.g. `nema_blit()`, `nema_fill_rect()`, `nema_fill_tri()` etc.).
19. **[M]** When using triangle culling, call `nema_tri_cull()` right before calling the rendering operation and call `nema_tri_cull(NEMA_CULL_NONE)` immediately afterwards. A misplaced `nema_tri_cull()` call can lead to some geometry PRIMs not to be drawn or to be drawn partially.
20. **[N]** Multiple PRIMs can be drawn within the same CL.
21. **[M]** For any drawing operation to take place, the assembled CL must be submitted for execution (`nema_cl_submit()`).
22. **[M]** Before deleting or reusing a CL, the GPU must have finished its execution (`nema_cl_wait()`).
23. **[N]** Once the CL is executed, it can be resubmitted (`nema_cl_submit()`), resetted and reused (`nema_cl_rewind()`), or destroyed (`nema_cl_destroyed()`).

2.2 Frame is different than expected

1. **[M]** FB parameters must be set properly (address, width, height, stride, color format).
2. **[N]** The correct FB color format should be used which should match the DC format.
3. **[M]** Clipping rectangle must be set within the boundaries of the FB (`nema_set_clip(x,y,w,h)`), and must fulfill the following conditions:
 - a. $x \geq 0$ and $y \geq 0$
 - b. $x+w \leq \text{fb.RESX}$ and $y+h \leq \text{fb.RESY}$

4. **[M]** The proper blending mode should be set depending on the desired rendering operation (fill or blit).
5. **[M]** In case of blitting, the source TEX containing the texture should be bound before rendering (`nema_bind_src()`).
6. **[M]** Source TEX parameters should be set correctly (address, width, height, stride, color format, wrapping mode).
7. **[M]** The appropriate rendering function should be used for the desired drawing operation (e.g. `nema_blit()`, `nema_fill_rect()`, `nema_blit_rect_fit()`).
8. **[M]** The CL must have finished its execution (`nema_cl_wait()`) before reusing or destroying it .
9. **[M]** TSc TEX and FB formats should be used only when available on the hardware implementation (GPU).
10. **[M]** Not all color formats are valid for FB. Please see NEMA| pico Software Manual for more details.
11. **[M]** When using depth, `nema_bind_depth_buffer()` and `nema_set_depth()` must be called before calling `nema_enable_depth()`.
12. **[A]** When using depth, call `nema_enable_depth(1)` right before calling the rendering operation and call `nema_enable_depth(0)` immediately afterwards. A misplaced `nema_enable_depth()` call can lead to various artifacts.
13. **[M]** When using gradient, `nema_set_gradient()` should be called before calling `nema_enable_gradient()`.
14. **[A]** When using gradient, call `nema_enable_gradient(1)` right before calling the rendering operation and call `nema_enable_gradient(0)` immediately afterwards. A misplaced `nema_enable_gradient()` call can lead to various artifacts.
15. **[M]** When using triangle culling, call `nema_tri_cull()` right before calling the rendering operation and call `nema_tri_cull(NEMA_CULL_NONE)` immediately afterwards. A misplaced `nema_tri_cull()` call can lead to some geometry PRIMs not to be drawn or to be drawn partially.
16. **[M]** Culling must be used only when rendering triangle PRIMs. Having culling enabled when rendering other PRIMs (rectangles, quadrilaterals) results to undefined behavior.

2.3 Memory Leaks

1. **[A]** Delete any allocated CL if not needed anymore. (`nema_cl_destroy()`)
2. **[A]** Delete any allocated BO if are not needed anymore (`nema_buffer_destroy()`)
3. **[A]** Avoid generic (not useful to the GPU) memory allocations within the graphics memory. Deallocate memory accordingly.

2.4 Performance Issues

1. **[A]** Identify the bottleneck (CPU, GPU, memory bandwidth etc) and optimize accordingly.

2. **[N]** A GPU is orders of magnitude faster than a CPU when rendering graphics. Use hardware graphics acceleration when possible to achieve a better performance.
3. **[A]** Offload the CPU by reusing assembled CL when possible. The same CL can be submitted multiple times for execution.
4. **[A]** Multiple CLs can be submitted at a time. In this case, just wait for the execution of the last submitted one.
5. **[A]** Use a pair of CLs for pipelining, prepare a new CL while the submitted one is being executed by the GPU.
6. **[A]** Draw only needed geometry to minimize the number of PRIMs.
7. **[A]** Adjust Clip Rectangle to the area to be drawn.
8. **[A]** Use triangle culling when drawing 3D geometries to minimize the number of PRIMs to draw.
9. **[A]** Use dirty regions to partially update the frame content - use clipping accordingly.
10. **[A]** Minimize overdraws. Overdraws areas can be inspected for debugging using the `nema_debug_overdraws()` call. This function disables gradients and textures and forces blending modes to `NEMA_BL_ADD`. By doing this, the user can inspect whether different regions are being drawn one of top of the other, detecting possible over-drawn regions that could have an impact on the performance.
11. **[N]** Excessive use of perspective transformations (`nema_matrix3x3`, `nema_matrix4x4`) can increase the CPU load significantly.
12. **[A]** Using appropriate graphics assets (for example, by minimizing texture resolution and selecting a suitable format) can reduce memory usage and bandwidth.
13. **[A]** In GPU-bound applications, prefer point-sampling instead of bilinear texture filtering, when the rendering quality is acceptable.

3 Breakpoints

In order to further debug applications, NEMA| GFX implements support for breakpoints. Breakpoints can be inserted at different positions within a CL in order to pause its execution. When the GPU reaches a breakpoint, it signals the CPU about it and pauses its execution. The GPU will continue its operation when the CPU instructs it so. While the GPU is paused, the user can have an overview of the state of the GPU (by calling respective functions) and the rendering process. In order to use the breakpoints, NEMA| GFX offers the following functions:

- `void nema_brk_enable(void)`: enables the breakpoints.
- `void nema_brk_disable(void)`: disables the breakpoints.
- `int nema_brk_add(void)`: adds a new breakpoint to the CL and returns its id (positive value).
- `int nema_brk_wait(int brk_id)`: the CPU waits for the GPU to reach the breakpoint with *id* equal to *brk_id*. Setting the *brk_id* = 0 will make the CPU wait until the GPU encounters the next breakpoint, regardless of its *id*. The function returns the *id* of the encountered breakpoint.
- `void nema_brk_continue(void)`: instruct the GPU to resume execution.

The following example illustrates the usage of breakpoints. The GPU is programmed to draw three rectangles filled with different colors (red, green and blue). A breakpoint is placed after each respective function call. The assembled CL is submitted to the GPU. The CPU will wait for the GPU to reach the the first breakpoint (brk1). While the GPU is paused, the id of the breakpoint along with the current color value are printed. Afterwards the CPU instructs the GPU to continue its operation. It then waits for the GPU to reach the third breakpoint (brk3). The second breakpoint is intentionally ignored. When brk3 is encountered, its id and the drawing color will be printed. The CPU instructs again the GPU to continue and the program finishes its execution.

```
void fill_rectangles(void){  
  
    nema_cmdlist_t cl = nema_cl_create();  
    nema_cl_bind(&cl);  
    nema_cl_branch(&context_cl);  
  
    //Enable breakpoints  
    nema_brk_enable();  
  
    nema_set_blend_fill( NEMA_BL_SRC );  
    nema_fill_rect(10, 10, 20, 20, nema_rgba(255, 0, 0, 255));  
    //Add first breakpoint  
    int brk1 = nema_brk_add();  
    nema_fill_rect(40, 10, 20, 20, nema_rgba(0, 255, 0, 255));  
    //Add second breakpoint  
    int brk2 = nema_brk_add();  
    nema_fill_rect(70, 10, 20, 20, nema_rgba(0, 0, 255, 255));  
    //Add third breakpoint
```

```
int brk3 = nema_brk_add();

nema_cl_unbind();
//Submit Command List for execution
nema_cl_submit(&cl);

int brk;

//Wait for the GPU to reach any breakpoint
brk = nema_brk_wait(0);
printf("Breakpoint %d: Texture color register value: %x\n", brk, nema_reg_read(NE▶
MA_DRAW_COLOR));
//Let the GPU resume execution
nema_brk_continue();

//Wait for the GPU to reach the third breakpoint (ignore any other breakpoints)
brk = nema_brk_wait(brk3);
printf("Breakpoint %d: Texture color register value: %x\n", brk, nema_reg_read(NE▶
MA_DRAW_COLOR));
//Let the GPU resume execution
nema_brk_continue();

//Wait for the Command List to finish execution
nema_cl_wait(&cl);
nema_cl_destroy(&cl);

//Disable breakpoints
nema_brk_disable();
}
```

4 GPU Status

NEMA| GFX is able to communicate with NEMA GPUs by writing or reading the GPU registers as well as by creating command lists and submitting them for execution to NEMA GPU. Writing invalid values in such structs that are read by the GPU (registers and command lists) could potentially make the GPU stall. For instance, when the framebuffer address has not been set at all, the GPU will attempt to write non accessible memory, and will stall. In this case (GPU stall), the necessary information about which subcomponent of the GPU encountered a malfunction can be found in the GPU status register (0x0fc). When the GPU is idle, it contains the zero value ('0'), otherwise it contains a value different than zero. This value denotes the GPU subcomponents that are currently busy. It can be decoded as follows:

Bit	Function
31	System busy
30	Memory System Unit busy
29	Command List bus busy
28	Command List Processor busy
27-24	Rasterizer Unit busy
23-13	Reserved
12	Render Output Unit busy
11-9	Reserved
8	Texture Map Unit busy
7-5	Reserved
4	Fragment Processor busy
3-1	Reserved
0	Fragment Processor Scheduler busy